

Dynamic Workflows with Airflow and Python

By Thejas Raju

PyLadies Dublin


Date: September 17, 2024



Thejas Raju

He/Him

Software Engineer
Pinterest



Agenda

- 1 Intro to Apache Airflow** [Go to section](#)
- 2 Workflows as Code** [Go to section](#)
- 3 Dynamic Workflows** [Go to section](#)
- 4 Best practices** [Go to section](#)
- 5 Conclusion** [Go to section](#)

**Apache Airflow is a platform
created by the community to
programmatically **author,**
schedule and monitor
workflows.**

Features

**Pure
python &
flexibility**

**Useful UI &
integrations**

**Ease of use
& open
source
community**

☒ Auto-refresh

DAG: example_bash_operator

Schedule: 0 0 * * *

Next Run ID: 2024-04-01, 20:00:00

04/02/2024, 11:55:37 AM

All Run Types

All Run States

Clear Filters

Auto-refresh

25

Press **shift** + **/** for Shortcuts **deferred** **failed** **queued** **removed** **restarting** **running** **scheduled** **shutdown** **skipped** **success** **up_for_reschedule** **up_for_retry** **upstream_failed** **no_status**DAG Run Task
example_bash_operator / 2024-04-01, 20:00:00 EDT / run_after_loop

Clear task

Mark state as...

Filter Tasks

Details

Graph

Gantt

Code

Audit Log

Logs

XCom

Task Duration

Layout:

Left -> Right

runme_2
success
BashOperatorthis_will_skip
skipped
BashOperatorrunme_1
success
BashOperatorrun_after_loop
success
BashOperatorrun_this_last
skipped
EmptyOperatorrunme_0
success
BashOperatoralso_run_this
success
BashOperator

React Flow

Spinner at Pinterest

**1.3k+ active
users**

**14k+ daily
DAG runs**

**131k+ daily
task
instances**

**160K+ daily
UI/API
requests**

Workflows as Code

The main characteristic of Airflow workflows is that all workflows are defined in Python code.

Workflows as Code

The main characteristic of Airflow workflows is that all workflows are defined in Python code.

Directed Acyclic Graphs (DAG)

A DAG is a collection of all the tasks you want to run, organized in a way that explicitly lays out their dependencies and execution order in a workflow.

Tasks, Operators, and Dependencies

Tasks are the individual units of work, operators define specific types of tasks, and dependencies are the relationships determining the order in which tasks run.

```
from airflow import DAG
from airflow.operators.dummy_operator import DummyOperator
from airflow.operators.python_operator import PythonOperator
from datetime import datetime

with DAG('sample_dag', start_date=datetime(2024, 9, 10)) as dag:

    start = DummyOperator(task_id='start')

    print_hello = PythonOperator(task_id='print_hello',
                                  python_callable=lambda: print('Hello PyLadies Dublin!'))

    print_welcome = PythonOperator(task_id='print_welcome',
                                     python_callable=lambda: print('Welcome to Pinterest!'))

    end = DummyOperator(task_id='end')

    start >> [print_hello, print_welcome] >> end
    start >> end
```

On DAG: sample_dag

schedule: 1 day, 0:00:00

Next Scheduled Execution Date 🕒 : 2024-09-13 00:00:00 UTC

Start Date of Next DagRun 🕒 : 2024-09-14 00:00:00 UTC

- Graph View
 - Grid View
 - DagRun History
 - Duration History
 - Task Tries
 - Landing Times
 - Gantt
 - Details
 - <> Code
 - Trigger DAG
- Dag Parsing Error Audit Log Refresh


Hologram and PCloud Token Info

success Base date: 2024-09-13 14:47:43+0 Number of runs: 25 Run: dev__traju__manual__2024-09-13T14:47:42.123140+00:00 Layout: Left->Right Go Search for...

DummyOperator PythonOperator

no_status upstream_failed queued up_for_retry up_for_reschedule skipped backfill_skipped trigger_skipped failed running success

☐ show critical path



```
[2024-09-12 13:25:12,662+00:00] [9417] {stage_base.py:190} INFO - Exiting stage ExecutorStage for task {'environment': 'devrestricted-traju', 'dag_id': 'sample_dag', 'task_id': 'print_hello', 'execution_date': '2024-09-11T00:00:00+00:00', 'try_number': 2}
[2024-09-12 13:25:12,679+00:00] [9417] {taskinstance.py:1058} INFO -
=====
[2024-09-12 13:25:12,680+00:00] [9417] {retry_restriction_utils.py:110} INFO - Auto retries is disabled because of either a manual retry or hitting the retry restriction threshold.
[2024-09-12 13:25:12,680+00:00] [9417] {retry_restriction_utils.py:121} INFO - For more information about retry restriction, please refer to: https://w.pinadmin.com/display/DataTeam/Retry+Restriction+User+Guide
[2024-09-12 13:25:12,680+00:00] [9417] {taskinstance.py:1073} INFO - Starting attempt 2 of 2
[2024-09-12 13:25:12,691+00:00] [9417] {taskinstance.py:1092} INFO - Executing <Task(PythonOperator): print_hello> on 2024-09-11T00:00:00+00:00
[2024-09-12 13:25:12,691+00:00] [9417] {base_task_runner.py:122} INFO - Running: ['airflow', 'tasks', 'run', 'sample_dag', 'print_hello', '2024-09-11T00:00:00+00:00', '--job_id', '5', '--pool', 'default_pool', '--raw', '-sd', 'DAGS_FOLDER/test_hello_world.py', '--cfg_path', '/tmp/tmpizqoze9m']
[2024-09-12 13:25:15,571+00:00] [9473] {task_command.py:394} INFO - Running <TaskInstance: sample_dag.print_hello 2024-09-11T00:00:00+00:00 [running]> on host devrestricted-traju.ec2.pin220.com
[2024-09-12 13:25:15,575+00:00] [9473] {taskinstance.py:1197} INFO - [JobExecutionRecord] Inserting new JobExecutionRecord into db
[2024-09-12 13:25:15,586+00:00] [9473] {taskinstance.py:1216} INFO - [JobExecutionRecord] Inserted JobExecutionRecord into db: JobExecutionRecord(workflow_name=sample_dag, workflow_cluster=devrestricted-traju, job_name=print_hello, instance_id=dev_traju_scheduled_2024-09-11T00:00:00+00:00 try_number=2, status=started, application_id_string=, jss_id_string=, start_time=2024-09-12 13:25:15.575895, end_time=None, project=ii, tier=tier3, data={"pinairflow_build": "9a9e13176cb3a97b2a9631e8f31ef63bb28b60e2", "pinboard_build": "4fb1d272eadaaf66495f61d21f7152bf4c6c18d7", "spinner_workflows_build": null})
[2024-09-12 13:25:15,586+00:00] [9473] {taskinstance.py:1260} INFO - Add JobExecutionRecord info in task instance context
[2024-09-12 13:25:15,590+00:00] [9473] {python_operator.py:135} INFO - Exporting the following env vars:
AIRFLOW_CTX_DAG_ID=sample_dag
AIRFLOW_CTX_TASK_ID=print_hello
AIRFLOW_CTX_EXECUTION_DATE=2024-09-11T00:00:00+00:00
AIRFLOW_CTX_DAG_RUN_ID=dev_traju_scheduled_2024-09-11T00:00:00+00:00
[2024-09-12 13:25:15,590+00:00] [9473] {logging_mixin.py:100} INFO - Hello PyLadies Dublin!
[2024-09-12 13:25:15,590+00:00] [9473] {python_operator.py:147} INFO - Done. Returned value was: None
[2024-09-12 13:25:15,593+00:00] [9473] {taskinstance.py:1353} INFO - [JobExecutionRecord] Updating final status of JobExecutionRecord in the db
[2024-09-12 13:25:15,601+00:00] [9473] {taskinstance.py:1374} INFO - [JobExecutionRecord] Updated final status of JobExecutionRecord in the db: JobExecutionRecord(workflow_name=sample_dag, workflow_cluster=devrestricted-traju, job_name=print_hello, instance_id=dev_traju_scheduled_2024-09-11T00:00:00+00:00 try_number=2, status=success, application_id_string=, jss_id_string=, start_time=2024-09-12 13:25:16, end_time=2024-09-12 13:25:15.595301, project=ii, tier=tier3, data={"pinairflow_build": "9a9e13176cb3a97b2a9631e8f31ef63bb28b60e2", "pinboard_build": "4fb1d272eadaaf66495f61d21f7152bf4c6c18d7", "spinner_workflows_build": null})
[2024-09-12 13:25:15,605+00:00] [9473] {stage_base.py:190} INFO - Exiting stage TaskExecStage for task {'environment': 'devrestricted-traju', 'dag_id': 'sample_dag', 'task_id': 'print_hello', 'execution_date': '2024-09-11T00:00:00+00:00', 'try_number': 2}
[2024-09-12 13:25:15,605+00:00] [9473] {stage_base.py:190} INFO - Exiting stage CallBackStage for task {'environment': 'devrestricted-traju', 'dag_id': 'sample_dag', 'task_id': 'print_hello', 'execution_date': '2024-09-11T00:00:00+00:00', 'try_number': 2}
```

Principles

Scalable

Elegant

Extensible

Dynamic

Dynamic Workflows

Workflows that can automatically adjust and scale based on varying inputs, conditions, or external triggers.

Input-based Dynamic Workflows

Using environment variables

```
deployment = os.environ.get("DEPLOYMENT",  
                             "PROD")  
  
if deployment == "PROD":  
    task = Operator(param="prod-param")  
elif deployment == "DEV":  
    task = Operator(param="dev-param")  
  
// DAG code
```

Example of an unsafe approach:

keeping both the dev and prod data in the same location puts you at risk of overwriting prod data. This behavior will be blocked in the future.

```
prod_database = 'default'
dev_database = os.environ.get('user')

table_name = 'very_important_table'

output_dev_location = f's3://    bucket-name    /project/{dev_database}/{table_name}'
output_prod_location = 's3://    bucket-name    /project/{prod_database}/{table_name}'

...
```


Use Case: S3ResourceSpec

Input-based Dynamic Workflows at Pinterest

S3ResourceSpec

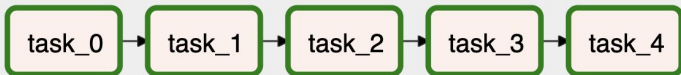
```
from airflow.lineage.pinterest.s3_resource
import S3ResourceSpec

# define the s3 url for the output table
output_s3_resource = S3ResourceSpec(
    base_url='s3://bucket-name/retention-30days/
    prefix_name/number_of_buckets',
    test_resource_retention_days=30,
)

output_s3_resource_entity =
output_s3_resource.get_resource_entity()
```

Condition-based Dynamic Workflows


Using iteration to generate tasks dynamically



```
from airflow import DAG
from airflow.operators.python_operator
import PythonOperator
from datetime import datetime

with DAG('sample_dag',
start_date=datetime(2024, 9, 10)) as dag:
    previous_task = None
    for i in range(5):
        task =
        PythonOperator(task_id=f'task_{i}',
python_callable=lambda: print(f'Hello from
task {i}'))
        if previous_task:
            previous_task >> task
        previous_task = task
```

Use Case: Data Profiling



```
def _push_tables_to_xcoms(context):
    resolved = []
    for row in context['job_output']:
        dataset_name = row[0]

        if dataset_name not in resolved:
            resolved.append(dataset_name)

    context['task_instance'].xcom_push(key='datasets',
value=resolved)

    context['task_instance'].xcom_push(key='wf_exec_date',
value=context['execution_date'].add(days=6))

    .strftime('%Y-%m-%d'))
```

```
def _set_config(context, dag_run_obj):
    """Set the dag_run.config to pass downstream."""
    dag_run_obj.payload = {
        'datasets':
context['task_instance'].xcom_pull(task_ids='get_tier_1
_candidates', key='datasets'),
        'wf_exec_date':
context['task_instance'].xcom_pull(task_ids='get_tier_1
_candidates', key='wf_exec_date')
    }
    LOG.info(dag_run_obj)
    LOG.info(dag_run_obj.payload)
    return dag_run_obj
```

Triggered DAG: ii_dp_daily

```
from  
airflow.providers.pinterest.models.dynamic_dag  
import DynamicDAG
```

```
dag = DynamicDAG(  
    dag_id='ii_data_profile',  
    schedule_interval=None,  
    default_args=default_args,  
    concurrency=20,  
    compute_layout=compute_layout,  
    max_active_runs=3,  
)
```

```
def compute_layout(dynamic_dag: DynamicDAG,  
                   execution_date: datetime = None,  
                   dagrun_conf: dict = None) -> None:  
  
    dag_name = dynamic_dag.dag_id  
  
    start = dynamic_dag.add_task_into_dynamic_dag(DummyOperator,  
    task_id='{}.start'.format(dag_name))  
  
    data_profiling_task = {}  
    populate_pincat_task = {}  
  
    if isinstance(dagrun_conf, dict) and "datasets" in dagrun_conf and  
    isinstance(dagrun_conf["datasets"], list):  
  
        exec_dt = dagrun_conf["wf_exec_date"]  
  
        for dataset_name in dagrun_conf["datasets"]:  
  
            start >> data_profiling_task[dataset_name] >>  
            populate_pincat_task[dataset_name]
```



Advantages

Flexibility

Scalability

Efficiency

Best Practices

Best Practices

**Optimize DAG
parsing**

**Manage
number of
tasks and
parallelism**

**Leverage
XComs
strategically**

Conclusion

Conclusion

- Airflow's capabilities provide extensive control over workflow automation and scheduling.
- Dynamically generating workflows can enhance adaptability and operational efficiency.
- Following best practices is essential for optimizing performance and maintainability.

Questions?

Thank you!



Thejas Raju
Data Engineering at Pinterest



www.linkedin.com/in/thejas-raju

Appendix

- [Spinner: Pinterest's Workflow Platform](#)
- [Airflow at Pinterest: Airflow Summit 2022](#)
- [Apache Airflow: Quick Start](#)